

How to evict file from system cache on Linux?

Asked 16 years, 2 months ago Modified 15 years, 5 months ago

Viewed 5k times



6



When running performance tests file system cache hit or miss can significantly influence test results. Therefore generally before running such tests used files are evicted from system cache. How to do that on Linux?

Clarification: If possible, the solution should not require root privileges.

linux

performance

testing

caching

file-io

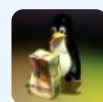
Share

edited Oct 2, 2008 at 13:28

Improve this question

Follow

asked Oct 2, 2008 at 11:53



Paweł Hajdan

18.5k ● 9 ● 52 ● 65

5 Answers

Sorted by:

Highest score (default)



As a superuser you can do the following:



10



To free pagecache:

- `echo 1 > /proc/sys/vm/drop_caches`

To free dentries and inodes:

- `echo 2 > /proc/sys/vm/drop_caches`

To free pagecache, dentries and inodes:

- `echo 3 > /proc/sys/vm/drop_caches`

This operation will not "lose" any data (caches are written out to disk before their data is dropped), however, to really make sure all cache is cleaned, you should sync first. E.g. all caches should be cleared if you run

```
sync; echo 3 > /proc/sys/vm/drop_caches
```

As I said, only a superuser (root) may do so.

[Share](#) [Improve this answer](#)

[Follow](#)

answered Oct 2, 2008 at 12:38



[Mecki](#)

133k ● 34 ● 259 ● 269



10



Ha, I have the answer:

```
#include <unistd.h>
#include <fcntl.h>
int main(int argc, char *argv[]) {
    int fd;
    fd = open(argv[1], O_RDONLY);
    fdatsync(fd);
}
```



```
posix_fadvise(fd, 0, 0, POSIX_FADV_DONTNEED);  
close(fd);  
return 0;  
}
```




This is from <http://insights.oetiker.ch/linux/fadvise.html>

Share Improve this answer

Follow

answered Oct 2, 2008 at 14:58



Paweł Hajdan

18.5k ● 9 ● 52 ● 65



There is a command line utility by Eric Wong that makes it easy to invoke `posix_fadvise`:

6

<http://git.bogomips.org/cgit/pcu.git/tree/README>



It's then as simple as



```
$ pcu-fadvise -a dontneed filename-to-evict
```



Share Improve this answer

Follow

answered Jul 22, 2009 at 20:12



user143174

61 ● 1 ● 1



Regarding use of `O_DIRECT`: that would perturb the results in another way. The kernel will attempt to DMA the filesystem data directly into your `read()` buffer, so it can be handed up to your application without any additional copy being done. Without `O_DIRECT` the kernel DMAs the file

0





data into the page cache, and copies it from the page cache to your read() buffer.



This is fine if your app is really going to use O_DIRECT in production. If you run performance tests with O_DIRECT and then remove O_DIRECT for production, your performance test will be unrealistic.

Share Improve this answer

answered Oct 2, 2008 at 13:55

Follow



DGentry

16.3k ● 8 ● 53 ● 66



-1

If you can put the test data in a separate filesystem then mounting the filesystem afresh for the test will give you empty caches.



If you list the test filesystem in /etc/fstab with the "user" option then you can mount it for the test without being superuser



Share Improve this answer

answered Oct 12, 2008 at 11:35

Follow



James

979 ● 2 ● 12 ● 22