



Profiles ▼

Maxim

Senior Software Engineer at Google



Member for 13 years, 8 months



Last seen this week



Zürich, Switzerland

Profile

Activity

Stats

53,728

reputation

4.2m

reached

731

answers

7

questions

About

Software Engineer at Google.

Badges

[View all badges](#)



27

gold badges



158

silver badges

● deep-learning

Jul 29, 2022

● machine-learning

Dec 20, 2018

● tensorflow

Jun 3, 2018

● conv-neural-network

Jul 27, 2020

● neural-network

Jul 21, 2018

● keras

Sep 18, 2019



211

bronze badges

● word2vec

Feb 19, 2019

● recurrent-neural-ne...

Oct 10, 2018

● lstm

Jul 2, 2018

Top tags

[View all tags](#)

tensorflow ●

2,790 score 480 posts 65 posts %

python ●

2,785 score 413 posts 56 posts %

machine-learning ●

2,119 score 383 posts 52 posts %

neural-network ●

1,106 score 200 posts 27 posts %

deep-learning ●

1,076 score 210 posts 28 posts %

keras ●

723 score 101 posts 14 posts %

Top posts

View all [questions](#), [answers](#), and [articles](#)

All

Questions

Answers

Articles

Score

Newest

A

968

Your CPU supports instructions that this TensorFlow binary was not compiled to use: AVX AVX2

Nov 10, 2017

A

151

How to choose cross-entropy loss in TensorFlow?

Oct 31, 2017

A

137

Instance Normalisation vs Batch normalisation

Jan 5, 2018

A

112

Batch Normalization in Convolutional Neural Network

Oct 11, 2017

Q

101

How to choose cross-entropy loss in TensorFlow?

Oct 31, 2017

A

87

What is the difference between a sigmoid followed by the cross entropy and `sigmoid_cross_entropy_with_logits` in TensorFlow?

Nov 11, 2017

A

81

What is a policy in reinforcement learning?

Sep 17, 2017

A

66

FutureWarning: Conversion of the second argument of `issubdtype` from `'float'` to `'np.floating'` is deprecated

Feb 13, 2018



59

What is right batch normalization function in Tensorflow?

Dec 28, 2017



56

What is the difference between xgb.train and xgb.XGBRegressor (or xgb.XGBClassifier)?

Nov 11, 2017

Top Meta posts

 2  0



7

Broken navigation in android app



3

Stack Overflow API ignores Accept-Encoding header

Top network posts

[View all network posts](#)



73

Do scientists know what is happening inside artificial neural networks?



63

How is softmax_cross_entropy_with_logits different from softmax_cross_entropy_with_logits_v2?



52

What's the difference between variance scaling initializer and xavier initializer?



43

Why are the weights of RNN/LSTM networks shared across time?



32

Why ReLU is better than the other activation functions

Collectives



NLP
2024 • Member

Communities

View all

 Stack Overflow	53.7k
 Cross Validated	3.3k
 Artificial Intelligence	2k
 Data Science	910
 Computer Science	640